

# Stat4620\_Project

Project Group 1

2024-11-20

```
library(ISLR)
library(pls)
```

```
##
## Attaching package: 'pls'

## The following object is masked from 'package:stats':
##
##   loadings
```

```
library(ggplot2)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v lubridate  1.9.2      v tibble    3.2.1
## v purrr      1.0.2      v tidyr     1.3.0
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x tidyr::expand() masks Matrix::expand()
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x tidyr::pack()    masks Matrix::pack()
## x tidyr::unpack() masks Matrix::unpack()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(broom)
library(dplyr)
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##     select
```

```
library(corrplot)
```

```
## corrplot 0.95 loaded
##
## Attaching package: 'corrplot'
##
## The following object is masked from 'package:pls':
##
##     corrplot
```

```
library(randomForest)
```

```
## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##     combine
##
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift
##
## The following object is masked from 'package:pls':
##
##     R2
```

```
library(gridExtra)
```

```
##
## Attaching package: 'gridExtra'
##
```

```
## The following object is masked from 'package:randomForest':
##
##      combine
##
## The following object is masked from 'package:dplyr':
##
##      combine
```

```
train_data = read.csv("train.csv")
test_data = read.csv("test_new.csv")
```

## #Part I: Exploratory Data Analysis

The AMES Housing data set contains information regarding to house prices and the characteristics of them. Variables ranges from numerical and categorical types of property locations, rooms and house furnishings.

```
# Check missing values for each column
missing_counts <- colSums(is.na(train_data))
missing_features <- missing_counts[missing_counts > 0]
missing_features
```

```
## LotFrontage      Alley  MasVnrType  MasVnrArea  BsmtQual  BsmtCond
##          259        1369           8           8          37          37
## BsmtExposure BsmtFinType1 BsmtFinType2  Electrical  FireplaceQu  GarageType
##          38           37           38           1          690          81
## GarageYrBlt GarageFinish  GarageQual  GarageCond      PoolQC      Fence
##          81           81           81           81        1453        1179
## MiscFeature
##          1406
```

There is one variable (LotFrontage) that contained a lot of actual missing values and thus we will drop it. We will also drop the ID column in the data set as it's used as an identifier and has no useful information. Upon analyzing the remaining missing features with NAs, we realized those NAs represent an actual category and are not missing data values, so we will keep them in the dataset for now.

```
train_data = train_data[, !(names(train_data) %in% c("Id", "LotFrontage"))]
```

We'll also drop categorical variables that don't provide a good split of the data space. Doing this will further simplify the number of features without losing any important patterns or information. Kaggle provides us a comprehensive view of the percentage break down of the buckets in the categorical variables. We'll drop variables that have buckets that exceed 85% of the observations.

```
train_data = train_data[, !(names(train_data) %in% c("Street", "Alley", "PoolQC", "MiscFeature", "LandC
```

We will then fill in the NAs for the remaining variables with missing values, replacing NAs in categorical variables with "None". There are two remaining continuous variables with missing values: GarageYrBuilt and MasVnrArea. For GarageYrBuilt, we will replace the NAs with the median value in that variable, but for "MasVnrArea", we will replace with the value 0 to correspond with the 8 missing values of categorical variable "MasVnrType".

```
summary(train_data)
```

```
##      MSSubClass      MSZoning      LotArea      LotShape
##  Min.   : 20.0    Length:1460    Min.    : 1300    Length:1460
## 1st Qu.: 20.0    Class :character    1st Qu.: 7554    Class :character
## Median : 50.0    Mode  :character    Median : 9478    Mode  :character
## Mean   : 56.9
## 3rd Qu.: 70.0
## Max.   :190.0
##
##      LotConfig      Neighborhood      BldgType      HouseStyle
## Length:1460      Length:1460      Length:1460      Length:1460
## Class :character    Class :character    Class :character    Class :character
## Mode  :character    Mode  :character    Mode  :character    Mode  :character
##
##
##
##      OverallQual      OverallCond      YearBuilt      YearRemodAdd
##  Min.   : 1.000    Min.    :1.000    Min.    :1872    Min.    :1950
## 1st Qu.: 5.000    1st Qu.:5.000    1st Qu.:1954    1st Qu.:1967
## Median : 6.000    Median :5.000    Median :1973    Median :1994
## Mean   : 6.099    Mean    :5.575    Mean    :1971    Mean    :1985
## 3rd Qu.: 7.000    3rd Qu.:6.000    3rd Qu.:2000    3rd Qu.:2004
## Max.   :10.000    Max.    :9.000    Max.    :2010    Max.    :2010
##
##      RoofStyle      Exterior1st      Exterior2nd      MasVnrType
## Length:1460      Length:1460      Length:1460      Length:1460
## Class :character    Class :character    Class :character    Class :character
## Mode  :character    Mode  :character    Mode  :character    Mode  :character
##
##
##
##      MasVnrArea      ExterQual      Foundation      BsmtQual
##  Min.   : 0.0    Length:1460      Length:1460      Length:1460
## 1st Qu.: 0.0    Class :character    Class :character    Class :character
## Median : 0.0    Mode  :character    Mode  :character    Mode  :character
## Mean   : 103.7
## 3rd Qu.: 166.0
## Max.   :1600.0
## NA's   :8
##      BsmtExposure      BsmtFinType1      BsmtFinSF1      BsmtFinSF2
## Length:1460      Length:1460      Min.   : 0.0    Min.   : 0.00
## Class :character    Class :character    1st Qu.: 0.0    1st Qu.: 0.00
## Mode  :character    Mode  :character    Median : 383.5    Median : 0.00
##
##      Mean : 443.6    Mean : 46.55
##      3rd Qu.: 712.2    3rd Qu.: 0.00
##      Max.   :5644.0    Max.   :1474.00
##
##      BsmtUnfSF      TotalBsmtSF      HeatingQC      X1stFlrSF
##  Min.   : 0.0    Min.    : 0.0    Length:1460      Min.   : 334
## 1st Qu.: 223.0    1st Qu.: 795.8    Class :character    1st Qu.: 882
```

```

## Median : 477.5    Median : 991.5    Mode :character    Median :1087
## Mean   : 567.2    Mean   :1057.4      Mean   :1163
## 3rd Qu.: 808.0    3rd Qu.:1298.2     3rd Qu.:1391
## Max.   :2336.0    Max.   :6110.0     Max.   :4692
##
##      X2ndFlrSF      LowQualFinSF      GrLivArea      BsmtFullBath
## Min.   : 0      Min.   : 0.000      Min.   : 334      Min.   :0.0000
## 1st Qu.: 0      1st Qu.: 0.000      1st Qu.:1130      1st Qu.:0.0000
## Median : 0      Median : 0.000      Median :1464      Median :0.0000
## Mean   : 347      Mean   : 5.845      Mean   :1515      Mean   :0.4253
## 3rd Qu.: 728      3rd Qu.: 0.000      3rd Qu.:1777      3rd Qu.:1.0000
## Max.   :2065      Max.   :572.000      Max.   :5642      Max.   :3.0000
##
##      BsmtHalfBath      FullBath      HalfBath      BedroomAbvGr
## Min.   :0.00000      Min.   :0.000      Min.   :0.0000      Min.   :0.000
## 1st Qu.:0.00000      1st Qu.:1.000      1st Qu.:0.0000      1st Qu.:2.000
## Median :0.00000      Median :2.000      Median :0.0000      Median :3.000
## Mean   :0.05753      Mean   :1.565      Mean   :0.3829      Mean   :2.866
## 3rd Qu.:0.00000      3rd Qu.:2.000      3rd Qu.:1.0000      3rd Qu.:3.000
## Max.   :2.00000      Max.   :3.000      Max.   :2.0000      Max.   :8.000
##
##      KitchenAbvGr      KitchenQual      TotRmsAbvGrd      Fireplaces
## Min.   :0.000      Length:1460      Min.   : 2.000      Min.   :0.000
## 1st Qu.:1.000      Class :character      1st Qu.: 5.000      1st Qu.:0.000
## Median :1.000      Mode  :character      Median : 6.000      Median :1.000
## Mean   :1.047                        Mean   : 6.518      Mean   :0.613
## 3rd Qu.:1.000                        3rd Qu.: 7.000      3rd Qu.:1.000
## Max.   :3.000                        Max.   :14.000      Max.   :3.000
##
##      FireplaceQu      GarageType      GarageYrBlt      GarageFinish
## Length:1460      Length:1460      Min.   :1900      Length:1460
## Class :character      Class :character      1st Qu.:1961      Class :character
## Mode  :character      Mode  :character      Median :1980      Mode  :character
##                               Mean   :1979
##                               3rd Qu.:2002
##                               Max.   :2010
##                               NA's   :81
##      GarageCars      GarageArea      WoodDeckSF      OpenPorchSF
## Min.   :0.000      Min.   : 0.0      Min.   : 0.00      Min.   : 0.00
## 1st Qu.:1.000      1st Qu.: 334.5      1st Qu.: 0.00      1st Qu.: 0.00
## Median :2.000      Median : 480.0      Median : 0.00      Median : 25.00
## Mean   :1.767      Mean   : 473.0      Mean   : 94.24      Mean   : 46.66
## 3rd Qu.:2.000      3rd Qu.: 576.0      3rd Qu.:168.00      3rd Qu.: 68.00
## Max.   :4.000      Max.   :1418.0      Max.   :857.00      Max.   :547.00
##
##      EnclosedPorch      X3SsnPorch      ScreenPorch      PoolArea
## Min.   : 0.00      Min.   : 0.00      Min.   : 0.00      Min.   : 0.000
## 1st Qu.: 0.00      1st Qu.: 0.00      1st Qu.: 0.00      1st Qu.: 0.000
## Median : 0.00      Median : 0.00      Median : 0.00      Median : 0.000
## Mean   : 21.95      Mean   : 3.41      Mean   : 15.06      Mean   : 2.759
## 3rd Qu.: 0.00      3rd Qu.: 0.00      3rd Qu.: 0.00      3rd Qu.: 0.000
## Max.   :552.00      Max.   :508.00      Max.   :480.00      Max.   :738.000
##
##      Fence      MiscVal      MoSold      YrSold

```

```
## Length:1460      Min.   :    0.00   Min.   : 1.000   Min.   :2006
## Class :character  1st Qu.:    0.00   1st Qu.: 5.000   1st Qu.:2007
## Mode :character  Median :    0.00   Median : 6.000   Median :2008
##                  Mean   :   43.49   Mean   : 6.322   Mean   :2008
##                  3rd Qu.:    0.00   3rd Qu.: 8.000   3rd Qu.:2009
##                  Max.   :15500.00   Max.   :12.000   Max.   :2010
##
## SaleCondition      SalePrice
## Length:1460      Min.   : 34900
## Class :character  1st Qu.:129975
## Mode :character  Median :163000
##                  Mean   :180921
##                  3rd Qu.:214000
##                  Max.   :755000
##
```

```
missing_counts <- colSums(is.na(train_data))
missing_features <- missing_counts[missing_counts > 0]
missing_features
```

```
## MasVnrType      MasVnrArea      BsmtQual BsmtExposure BsmtFinType1 FireplaceQu
##           8           8           37           38           37           690
## GarageType      GarageYrBlt      GarageFinish      Fence
##           81           81           81           1179
```

```
median_value <- median(train_data$GarageYrBlt, na.rm = TRUE)
train_data$GarageYrBlt[is.na(train_data$GarageYrBlt)] <- median_value
train_data$MasVnrArea[is.na(train_data$MasVnrArea)] <- 0

train_data[is.na(train_data)] <- "None"

colSums(is.na(train_data))#there are now no NA's
```

```
## MSSubClass      MSZoning      LotArea      LotShape      LotConfig
##           0           0           0           0           0
## Neighborhood      BldgType      HouseStyle      OverallQual      OverallCond
##           0           0           0           0           0
## YearBuilt      YearRemodAdd      RoofStyle      Exterior1st      Exterior2nd
##           0           0           0           0           0
## MasVnrType      MasVnrArea      ExterQual      Foundation      BsmtQual
##           0           0           0           0           0
## BsmtExposure      BsmtFinType1      BsmtFinSF1      BsmtFinSF2      BsmtUnfSF
##           0           0           0           0           0
## TotalBsmtSF      HeatingQC      X1stFlrSF      X2ndFlrSF      LowQualFinSF
##           0           0           0           0           0
## GrLivArea      BsmtFullBath      BsmtHalfBath      FullBath      HalfBath
##           0           0           0           0           0
## BedroomAbvGr      KitchenAbvGr      KitchenQual      TotRmsAbvGrd      Fireplaces
##           0           0           0           0           0
## FireplaceQu      GarageType      GarageYrBlt      GarageFinish      GarageCars
##           0           0           0           0           0
## GarageArea      WoodDeckSF      OpenPorchSF      EnclosedPorch      X3SsnPorch
##           0           0           0           0           0
```

```
##   ScreenPorch      PoolArea      Fence      MiscVal      MoSold
##           0           0           0           0           0
##      YrSold SaleCondition      SalePrice
##           0           0           0
```

Now we will look at the correlation matrix of our continuous predictors and address any irrelevant features.

```
# Correlation matrix for numeric features
train_data_numeric <- train_data[sapply(train_data, is.numeric)]
cor_matrix <- cor(train_data_numeric)

#subset(as.data.frame.table(cor_matrix), abs(Freq) < 1 & abs(Freq) > 0.75)

cor_sal <- cor_matrix[, "SalePrice"]
cor_sal
```

```
##   MSSubClass      LotArea OverallQual OverallCond      YearBuilt
##   -0.08428414  0.26384335  0.79098160 -0.07785589  0.52289733
##   YearRemodAdd  MasVnrArea BsmtFinSF1  BsmtFinSF2  BsmtUnfSF
##   0.50710097  0.47261450  0.38641981 -0.01137812  0.21447911
##   TotalBsmtSF   X1stFlrSF  X2ndFlrSF  LowQualFinSF  GrLivArea
##   0.61358055  0.60585218  0.31933380 -0.02560613  0.70862448
##   BsmtFullBath  BsmtHalfBath  FullBath  HalfBath  BedroomAbvGr
##   0.22712223  -0.01684415  0.56066376  0.28410768  0.16821315
##   KitchenAbvGr  TotRmsAbvGrd  Fireplaces  GarageYrBlt  GarageCars
##   -0.13590737  0.53372316  0.46692884  0.46675365  0.64040920
##   GarageArea    WoodDeckSF  OpenPorchSF  EnclosedPorch  X3SsnPorch
##   0.62343144  0.32441344  0.31585623  -0.12857796  0.04458367
##   ScreenPorch   PoolArea    MiscVal    MoSold    YrSold
##   0.11144657  0.09240355  -0.02118958  0.04643225  -0.02892259
##   SalePrice
##   1.00000000
```

```
# All variables not highly correlated with SalePrice
names(cor_sal[abs(cor_sal) < 0.5])
```


```
## [1] "MSSubClass" "LotArea" "OverallCond" "MasVnrArea"
## [5] "BsmtFinSF1" "BsmtFinSF2" "BsmtUnfSF" "X2ndFlrSF"
## [9] "LowQualFinSF" "BsmtFullBath" "BsmtHalfBath" "HalfBath"
## [13] "BedroomAbvGr" "KitchenAbvGr" "Fireplaces" "GarageYrBlt"
## [17] "WoodDeckSF" "OpenPorchSF" "EnclosedPorch" "X3SsnPorch"
## [21] "ScreenPorch" "PoolArea" "MiscVal" "MoSold"
## [25] "YrSold"
```

We will remove all the continuous variables that are not highly correlated with our response variable, SalePrice, based on the correlation matrix above. Those continuous variables with a correlation value higher than 0.5 or lower than -0.5 will remain in our dataset.

```
train_data = train_data[, !(names(train_data) %in% c("MSSubClass", "LotArea", "OverallCond", "BsmtFinSF1", "BsmtFinSF2", "BsmtUnfSF", "X2ndFlrSF", "LowQualFinSF", "BsmtFullBath", "BsmtHalfBath", "HalfBath", "BedroomAbvGr", "KitchenAbvGr", "Fireplaces", "GarageYrBlt", "GarageCars", "WoodDeckSF", "OpenPorchSF", "EnclosedPorch", "X3SsnPorch", "ScreenPorch", "PoolArea", "MiscVal", "MoSold", "YrSold"))]
```

We will now look at the correlation between all predictor variables to see if there are any two that are highly correlated. If two of them are highly correlated then we will remove the one that is least correlated with the response variable.

```
# Create correlation plot
corrplot(cor_matrix, method = "color", tl.cex = 0.5)
```



```
# Print all relationships with 0.75 correlation or more
subset(as.data.frame(table(cor_matrix)), abs(Freq) < 1 & abs(Freq) > 0.75)
```

##	Var1	Var2	Freq
## 108	SalePrice	OverallQual	0.7909816
## 168	GarageYrBlt	YearBuilt	0.7771818
## 372	X1stFlrSF	TotalBsmtSF	0.8195300
## 407	TotalBsmtSF	X1stFlrSF	0.8195300
## 526	TotRmsAbvGrd	GrLivArea	0.8254894
## 771	GrLivArea	TotRmsAbvGrd	0.8254894
## 833	YearBuilt	GarageYrBlt	0.7771818
## 890	GarageArea	GarageCars	0.8824754
## 925	GarageCars	GarageArea	0.8824754
## 1263	OverallQual	SalePrice	0.7909816

In the table above we can see that 4 of the predictor variables are highly correlated with another 4 variables so we will remove those, keeping the ones with higher correlation to the response.

```
# Remove variables due to multicollinearity
train_data = train_data[, !(names(train_data) %in% c("GarageYrBlt", "X1stFlrSF", "TotRmsAbvGrd", "Garag
```



We will now look at all the categorical variables to see if they all have a unique distribution of SalePrice across different categories, deeming them useful.


```
# List of all categorical variables
categorical_vars <- c("MSZoning", "LotShape", "LotConfig", "Neighborhood", "BldgType", "HouseStyle", "R")



plot_list <- list()



# Loop through categorical variables and store plots in the list
for (var in categorical_vars) {
  x <- ggplot(train_data, aes_string(x = var, y = "SalePrice")) +
    geom_bar(stat = "summary", fun = "mean", fill = "steelblue") +
    labs(title = paste("SalePrice by", var),
         x = var, y = "Average SalePrice") +
    theme_minimal() +
    theme(axis.text.x = element_text(angle = 45, hjust = 1))

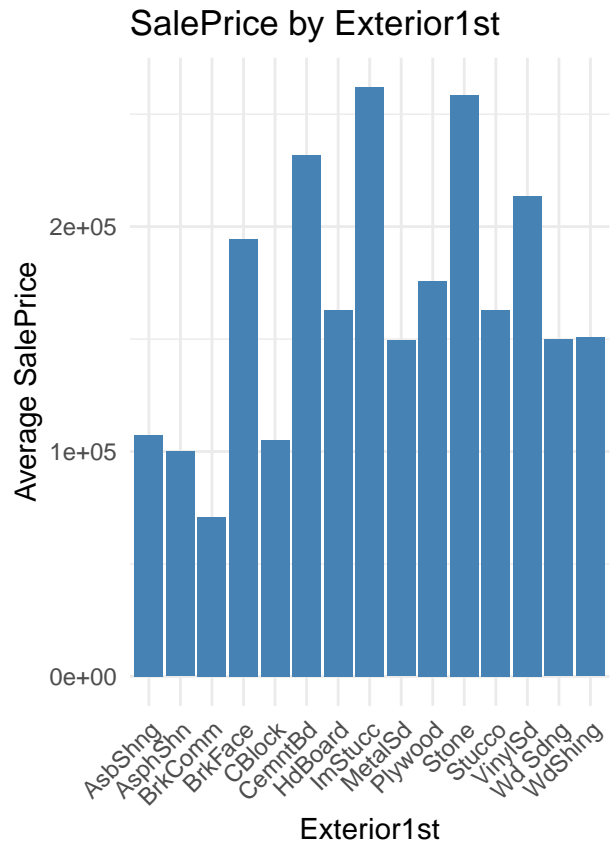

  # Add the plot to the list
  plot_list[[length(plot_list) + 1]] <- x
}

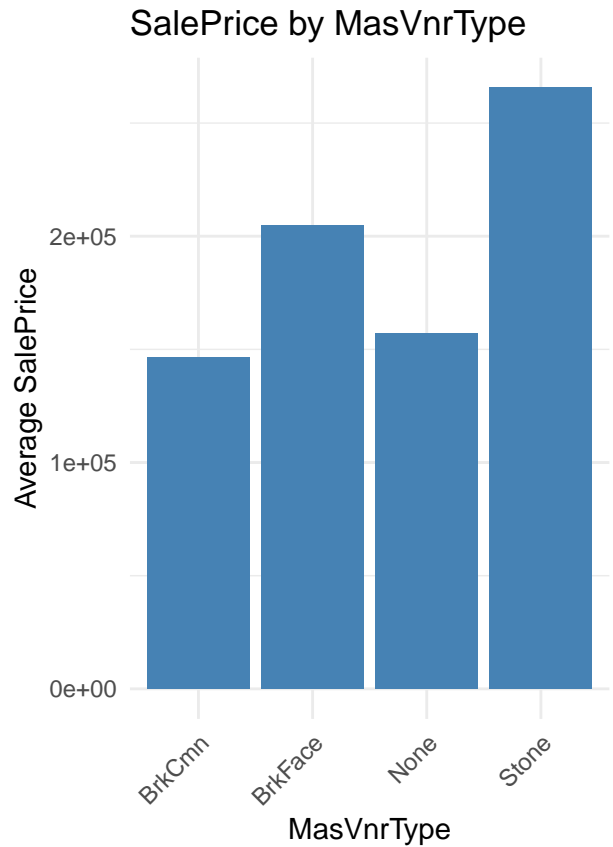

# Arrange and print the plots two at a time
for (i in seq(1, length(plot_list), by = 2)) {
  plots_to_print <- plot_list[i:min(i + 1, length(plot_list))]
  grid.arrange(grobs = plots_to_print, ncol = 2)
}
```




















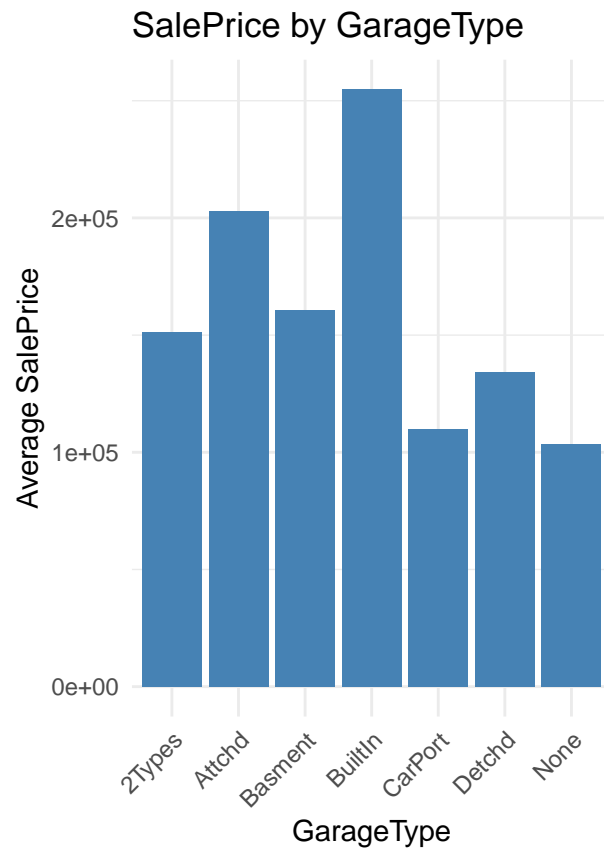


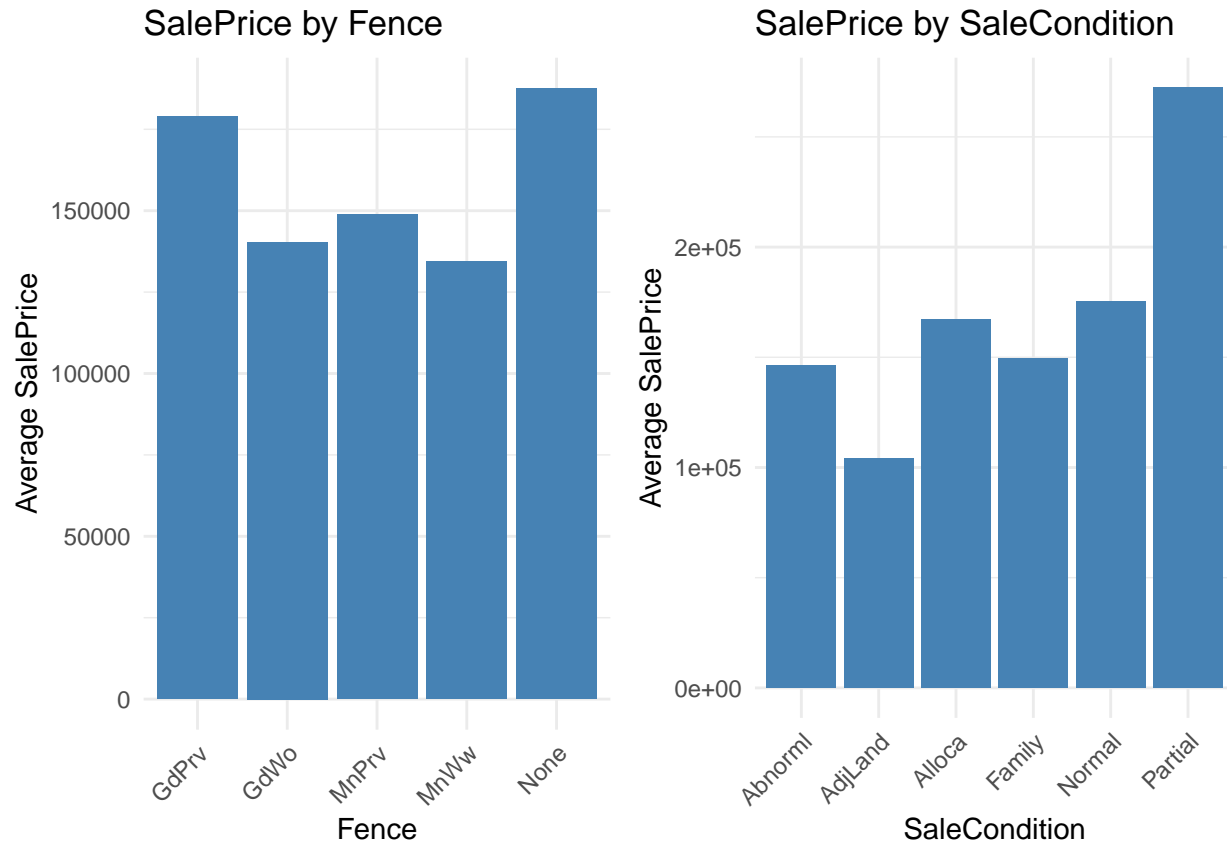












We can see that for each categorical variable that the SalePrice is different across each category in each categorical variable which is good and tells us that they will all be useful.

#### #Part II: Model Analysis

+After cleaning our data and performing EDA we are going to fit a regression tree model to our data.

+A regression tree model is a decision tree that predicts a continuous variable. It predicts by recursively partitioning the predictor space into smaller and smaller subregions the more we split the tree. The tree defines local regions using a step-function-like approach.

+The regression tree model would give us the best results because they work with complex data sets where there's a mix of categorical and numerical variables. Trees also don't have the typical linear regression interpretation, so we wouldn't need to create indicator variables to represent the categorical features. Lastly, regression trees allow us to easily interpret the results.

+The regression tree model makes minimal assumptions on the relationships in the data set. The assumption generally being that the data can be partitioned into subsets and that each split is independent and interpretable.

```
cv <- trainControl(method = "cv", number = 5) # 5-fold cross-validation
mtry_grid <- expand.grid(.mtry = c(15, 20, 25, 30)) # Tuning grid for mtry

# Train the model using random forest with cross-validation
set.seed(123)
rf_cv_model <- train(SalePrice ~ ., data = train_data, method = "rf", trControl = cv, tuneGrid = mtry_grid,
print(rf_cv_model)
```

## Random Forest

```
##
## 1460 samples
## 30 predictor
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 1169, 1169, 1167, 1168, 1167
## Resampling results across tuning parameters:
##
##  mtry  RMSE      Rsquared  MAE
##  15    30690.20  0.8621157  18241.05
##  20    29802.93  0.8685531  17814.46
##  25    29474.77  0.8697370  17700.93
##  30    29266.70  0.8710967  17554.12
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was mtry = 30.
```


After running cross-validation to select the best value for mtry which represents the number of predictors sampled for splitting at each node. We can see that mtry=30 gives the best results. We will now fit the model with 500 trees and mtry=30

```
set.seed(123)
rf_model <- randomForest(SalePrice ~ ., data = train_data, ntree = 500, mtry = 30, importance = TRUE)
print(rf_model)
```

```
##
## Call:
## randomForest(formula = SalePrice ~ ., data = train_data, ntree = 500, mtry = 30, importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 30
##
##              Mean of squared residuals: 934607442
##              % Var explained: 85.18
```

After fitting the model we got an  $R^2$  value of 0.8218, indicating that 82% of the variance in the response variable is explained by the selected features. This suggests that the model provided a strong fit for the data.

```
par(cex = 0.7)
varImpPlot(rf_model, type = 1) # Plot variable importance
```



We can see that the most important variable is **GrLiveArea** (Above ground living area square feet) which makes sense because larger houses will cost more. Along with that variable, **OverallQual** (Overall material and finish of the house) and **TotalBsmtSF** (Total square feet of the basement) are also very important to the model and separate themselves from the other variables. **TotalBsmtSF** is very similar to **GrLiveArea** and probably gives a similar value so that explains why it is so important and if the quality of the house is low then the price will also be lower.

We will now run the test data through the pre-processing and then evaluate its performance with the model.

```
#Pre-Processing on test_data
```

```
test_data = test_data[, !(names(test_data) %in% c("Id", "LotFrontage", "Street", "Alley", "PoolQC", "Mi
```

```
test_data$MasVnrArea[is.na(test_data$MasVnrArea)] <- 0
```

```
test_data[is.na(test_data)] <- "None"
```

```
test_x = test_data[, !(names(test_data) == "SalePrice")] #predictors of test data
```

```
test_y = test_data[, (names(test_data) == "SalePrice")] #response of test data
```

```
predictions <- predict(rf_model, newdata = test_x) #predict sale price on test data
```


```
#data frame of results
```

```
results = data.frame(Actual = test_y, Predicted = predictions)
```

```
# Scatterplot of actual vs predicted values
```

```
ggplot(data = results, aes(x = Actual, y = Predicted)) +  
  geom_point(color = "blue", alpha = 0.6) +
```

```
geom_abline(slope = 1, intercept = 0, color = "red", linetype = "solid") +
labs(title = "Actual vs Predicted Prices",
      x = "Actual Price",
      y = "Predicted Price") +
theme_minimal()
```



```
rmse <- sqrt(mean((predictions - test_y)^2)) #get rmse of predictions
cat("RMSE: ", rmse, "\n")
```

```
## RMSE: 26306.08
```

After running the test data through the model we can see that on average we are \$26,306.08 off from the actual sale price.